

Car Price Prediction

Dataset source: <https://archive.ics.uci.edu/ml/datasets/Automobile>

```
link='https://drive.google.com/uc?id=1UO
import pandas as pd
df=pd.read_csv(link)
df.head()
```

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody
0	1	3	alfa-romero giulia	gas	std	two	convertible
1	2	3	alfa-romero stelvio	gas	std	two	convertible
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback
3	4	2	audi 100 ls	gas	std	four	sedan
4	5	2	audi 100ls	gas	std	four	sedan

5 rows × 26 columns

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   car_ID          205 non-null    int64  
 1   symboling       205 non-null    int64  
 2   CarName         205 non-null    object  
 3   fueltype        205 non-null    object  
 4   aspiration      205 non-null    object  
 5   doornumber      205 non-null    object  
 6   carbody         205 non-null    object  
 7   drivewheel      205 non-null    object  
 8   enginelocation   205 non-null    object  
 9   wheelbase        205 non-null    float64 
 10  carlength       205 non-null    float64 
 11  carwidth        205 non-null    float64 
 12  carheight       205 non-null    float64 
 13  curbweight      205 non-null    int64  
 14  enginetype       205 non-null    object  
 15  cylindernumber   205 non-null    object  
 16  enginesize       205 non-null    int64  
 17  fuelsystem       205 non-null    object  
 18  boreratio        205 non-null    float64
```

```
19 stroke          205 non-null    float64
20 compressionratio 205 non-null    float64
21 horsepower       205 non-null    int64
22 peakrpm          205 non-null    int64
23 citympg          205 non-null    int64
24 highwaympg        205 non-null    int64
25 price            205 non-null    float64
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB
```

Cars are initially assigned a risk factor symbol associated with its price. Then, if it is more risky (or less), this symbol is adjusted by moving it up (or down) the scale. This process is called **symboling**. A value of +3 indicates that the auto is risky, -3 that it is probably pretty safe.

```
df[ 'CarName' ].head(25)
```

	CarName
0	alfa-romero giulia
1	alfa-romero stelvio
2	alfa-romero Quadrifoglio
3	audi 100 ls
4	audi 100ls
5	audi fox
6	audi 100ls
7	audi 5000
8	audi 4000
9	audi 5000s (diesel)
10	bmw 320i
11	bmw 320i
12	bmw x1
13	bmw x3
14	bmw z4
15	bmw x4
16	bmw x5
17	bmw x3
18	chevrolet impala
19	chevrolet monte carlo
20	chevrolet vega 2300
21	dodge rampage
22	dodge challenger se
23	dodge d200
24	dodge monaco (sw)

dtype: object

```
cars_name= pd.Series([i.split()[0] for i in df['CarName']])  
cars_name
```

```
0  
0    alfa-romero  
1    alfa-romero  
2    alfa-romero  
3        audi  
4        audi  
...  
200      volvo  
201      volvo  
202      volvo  
203      volvo  
204      volvo
```

205 rows × 1 columns

dtype: object

```
df['CarCompany']=cars_name  
df.drop(columns= ['car_ID','CarName'],inplace=True)  
df.head()
```

	symboling	fueltype	aspiration	doornumber	carbody	drivewheel	engine
0	3	gas	std	two	convertible	rwd	
1	3	gas	std	two	convertible	rwd	
2	1	gas	std	two	hatchback	rwd	
3	2	gas	std	four	sedan	fwd	
4	2	gas	std	four	sedan	4wd	

5 rows × 25 columns

```
df['CarCompany'].value_counts()
```

CarCompany	count
toyota	31
nissan	17
mazda	15
honda	13
mitsubishi	13
subaru	12
volvo	11
peugeot	11
dodge	9
volkswagen	9
buick	8
bmw	8
audi	7
plymouth	7
saab	6
porsche	4
isuzu	4
alfa-romero	3
jaguar	3
chevrolet	3
renault	2
maxda	2
vw	2
mercury	1
porcshce	1
Nissan	1
toyoutua	1
vokswagen	1

dtype: int64

```
df.loc[(df['CarCompany'] == "vw") | (df['CarCompany'] == "vokswagen"), 'CarC  
# porsche  
df.loc[(df['CarCompany']=='porcshce'), 'CarCompany'] = "porsche"  
# toyota  
df.loc[(df['CarCompany']=='toyouta'), 'CarCompany'] = "toyota"  
# nissan  
df.loc[(df['CarCompany']=='Nissan'), 'CarCompany'] = "nissan"  
# mazda  
df.loc[(df['CarCompany']=='maxda'), 'CarCompany'] = "mazda"  
  
df['CarCompany'].value_counts()
```

count

CarCompany

toyota	32
nissan	18
mazda	17
mitsubishi	13
honda	13
subaru	12
volkswagen	12
volvo	11
peugeot	11
dodge	9
buick	8
bmw	8
audi	7
plymouth	7
saab	6
porsche	5
isuzu	4
alfa-romero	3
chevrolet	3
jaguar	3
renault	2
mercury	1

dtype: int64

```
df.dtypes
```

	0
symboling	int64
fueltype	object
aspiration	object
doornumber	object
carbody	object
drivewheel	object
enginelocation	object
wheelbase	float64
carlength	float64
carwidth	float64
carheight	float64
curbweight	int64
enginetype	object
cylindernumber	object
enginesize	int64
fuelsystem	object
boreratio	float64
stroke	float64
compressionratio	float64
horsepower	int64
peakrpm	int64
citympg	int64
highwaympg	int64
price	float64
CarCompany	object

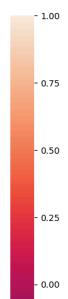
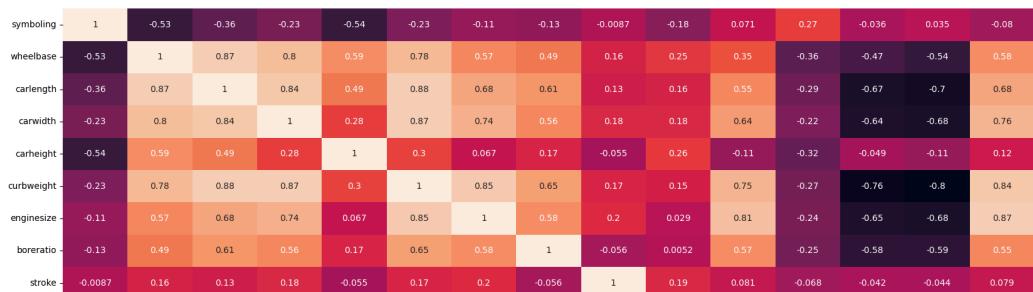
dtype: object

```
num_df= df.select_dtypes(include=['int','float'])
num_df.head()
```

	symboling	wheelbase	carlength	carwidth	carheight	curbweight	enginesize
0	3	88.6	168.8	64.1	48.8	2548	164.0
1	3	88.6	168.8	64.1	48.8	2548	164.0
2	1	94.5	171.2	65.5	52.4	2823	164.0
3	2	99.8	176.6	66.2	54.3	2337	164.0
4	2	99.4	176.6	66.4	54.3	2824	164.0

```
import matplotlib.pyplot as plt
import seaborn as sns
plt.figure(figsize=(25,10))
sns.heatmap(num_df.corr(), annot=True)
```

<Axes: >



```
df['doornumber'].value_counts()
```

```
count
```

```
doornumber
```

four	115
two	90

```
dtype: int64
```

```
df['cylindernumber'].value_counts()
```

```
count
```

```
cylindernumber
```

four	159
six	24
five	11
eight	5
two	4
twelve	1
three	1

```
dtype: int64
```

```
dict_words= {"two": 2, "three": 3, "four": 4, "five": 5, "six": 6, "eight": 8}
df['doornumber']= df['doornumber'].map(dict_words)
df['cylindernumber']=df['cylindernumber'].map(dict_words)
df.head()
```

	symboling	fueltype	aspiration	doornumber	carbody	drivewheel	engine
0	3	gas	std	2	convertible	rwd	
1	3	gas	std	2	convertible	rwd	
2	1	gas	std	2	hatchback	rwd	
3	2	gas	std	4	sedan	fwd	
4	2	gas	std	4	sedan	4wd	

5 rows × 25 columns

```
obj_df= df.select_dtypes(include='object')
obj_df
```

	fueltype	aspiration	carbody	drivewheel	enginelocation	enginetype
0	gas	std	convertible	rwd	front	dohc
1	gas	std	convertible	rwd	front	dohc
2	gas	std	hatchback	rwd	front	ohcv
3	gas	std	sedan	fwd	front	ohc
4	gas	std	sedan	4wd	front	ohc
...
200	gas	std	sedan	rwd	front	ohc
201	gas	turbo	sedan	rwd	front	ohc
202	gas	std	sedan	rwd	front	ohcv
203	diesel	turbo	sedan	rwd	front	ohc
204	gas	turbo	sedan	rwd	front	ohc

205 rows × 8 columns

Choosing between **one-hot encoding** and **label encoding** depends on the nature of the categorical variable you're working with and the machine learning algorithm you plan to use.

▼ 1. Nature of the Categorical Variable:

- **Ordinal Categorical Variables:**

- These are categories that have a natural order or ranking (e.g., "low," "medium," "high").
- **Use Label Encoding** because the integer values assigned by label encoding can reflect the ordinal relationship between the categories.
 - Example: "low" → 1, "medium" → 2, "high" → 3.

- **Nominal Categorical Variables:**

- These are categories that do not have an inherent order or ranking (e.g., "red," "blue," "green").
- **Use One-Hot Encoding** because label encoding might incorrectly imply a ranking or relationship between the categories that does not exist.
 - Example: "red" → [1, 0, 0], "blue" → [0, 1, 0], "green" → [0, 0, 1].

2. Machine Learning Algorithm:

- **Tree-Based Algorithms (e.g., Decision Trees, Random Forests):**

- These algorithms are not sensitive to the numerical nature of the input and can often handle label-encoded data effectively.
- **Label Encoding** can be used, especially for ordinal data.
- **Linear Algorithms (e.g., Logistic Regression, Linear Regression, SVM):**
 - These algorithms assume a linear relationship and can misinterpret label-encoded values as ordinal, which can introduce bias if the data is nominal.
 - **One-Hot Encoding** is generally preferred to prevent this issue and ensure the model does not impose a false ordinal relationship.

3. Number of Categories:

- **Few Categories (e.g., < 10):**
 - **One-Hot Encoding** is manageable even with a small number of categories, and it provides clear separations between categories.
- **Many Categories (e.g., > 10):**
 - **Label Encoding** may be preferred if there are a large number of categories, as one-hot encoding would create a very high-dimensional feature space, which can lead to the "curse of dimensionality" and increased computational cost.
 - **One-Hot Encoding** can still be used, but it might require dimensionality reduction techniques afterward.

4. Risk of Introducing Bias:

- **One-Hot Encoding** is safer in avoiding bias since it treats all categories equally without implying any order or relationship.
- **Label Encoding** can introduce bias if the algorithm assumes a relationship between the encoded values.

5. Data Sparsity:

- **One-Hot Encoding** results in sparse data (many zeros), which can be computationally inefficient for large datasets with many categories. Some models, however, are optimized to handle sparse data efficiently.
- **Label Encoding** results in dense data (single column), which is more computationally efficient.

```
# One hot encoding
df['carbody'].unique()

array(['convertible', 'hatchback', 'sedan', 'wagon', 'hardtop'],
      dtype=object)
```

```
pd.get_dummies(df['carbody'], dtype=int)
```

	convertible	hardtop	hatchback	sedan	wagon
0	1	0	0	0	0
1	1	0	0	0	0
2	0	0	1	0	0
3	0	0	0	1	0
4	0	0	0	1	0
...
200	0	0	0	1	0
201	0	0	0	1	0
202	0	0	0	1	0
203	0	0	0	1	0
204	0	0	0	1	0

205 rows × 5 columns

```
pd.get_dummies(df[['carbody', 'fueltype', 'aspiration']], dtype=int, drop_first=
```

	carbody_hardtop	carbody_hatchback	carbody_sedan	carbody_wagon	fuelty
0	0	0	0	0	0
1	0	0	0	0	0
2	0	1	0	0	0
3	0	0	0	1	0
4	0	0	0	1	0
...
200	0	0	0	1	0
201	0	0	0	1	0
202	0	0	0	1	0
203	0	0	0	1	0
204	0	0	0	1	0

205 rows × 6 columns

```
dummy_df= pd.get_dummies(obj_df, dtype=int, drop_first=True)
dummy_df
```

	<code>fueltype_gas</code>	<code>aspiration_turbo</code>	<code>carbody_hardtop</code>	<code>carbody_hatchback</code>	<code>carb</code>
0	1	0	0	0	0
1	1	0	0	0	0
2	1	0	0	0	1
3	1	0	0	0	0
4	1	0	0	0	0
...
200	1	0	0	0	0
201	1	1	0	0	0
202	1	0	0	0	0
203	0	1	0	0	0
204	1	1	0	0	0

205 rows × 43 columns

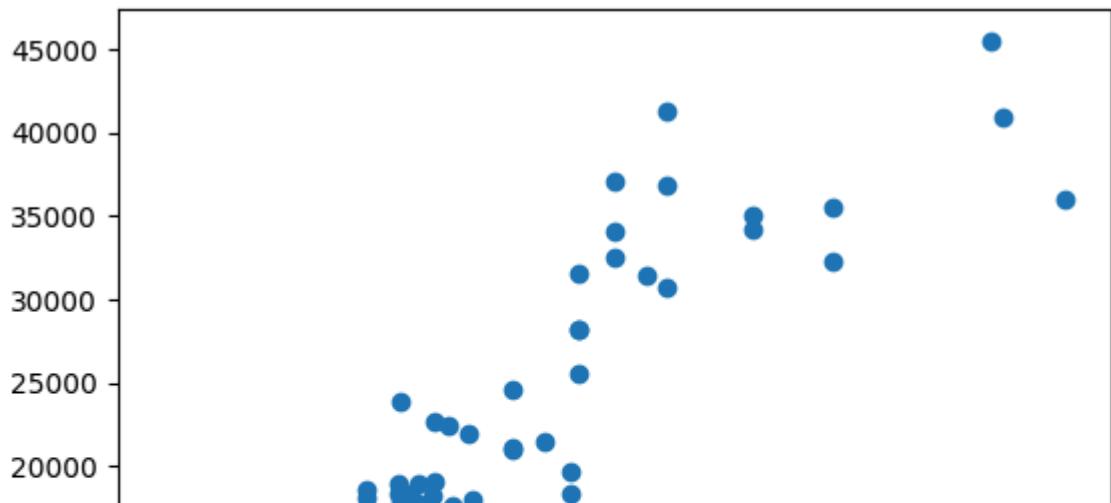
```
df.drop(columns=obj_df.columns,inplace=True)
df_new= pd.concat([df,dummy_df],axis=1)
df_new
```

	<code>symboling</code>	<code>doornumber</code>	<code>wheelbase</code>	<code>carlength</code>	<code>carwidth</code>	<code>carheight</code>	<code>curbwe</code>
0	3	2	88.6	168.8	64.1	48.8	1740
1	3	2	88.6	168.8	64.1	48.8	1740
2	1	2	94.5	171.2	65.5	52.4	1740
3	2	4	99.8	176.6	66.2	54.3	1740
4	2	4	99.4	176.6	66.4	54.3	1740
...
200	-1	4	109.1	188.8	68.9	55.5	1740
201	-1	4	109.1	188.8	68.8	55.5	1740
202	-1	4	109.1	188.8	68.9	55.5	1740
203	-1	4	109.1	188.8	68.9	55.5	1740
204	-1	4	109.1	188.8	68.9	55.5	1740

205 rows × 60 columns

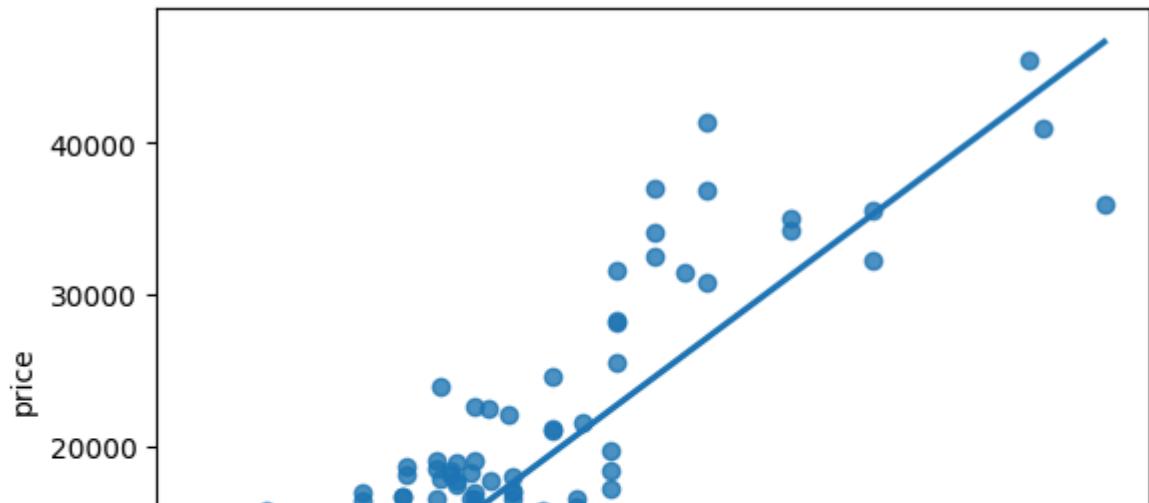
```
plt.scatter(df['enginesize'],df['price'])
```

```
<matplotlib.collections.PathCollection at 0x794c629ae6c0>
```



```
sns.regplot(x=df['enginesize'],y=df['price'],ci=None)
```

```
<Axes: xlabel='enginesize', ylabel='price'>
```



Simple Linear Regression

In simple linear regression, our job is to find this straight line which is called the **best fit line**. So we need to find the best fit line that can fit the most points in the scatter plot between the car price and the enginesize.

Let the equation of the best fit line be

$$y = mx + c$$

Here,

- y represents the price values on the y -axis
- x represents the enginesize values on the x -axis
- m is the slope of the line
- c is the intercept made by the line on the y -axis

The above equation can also be written as

$$\text{price} = m \times \text{enginesize} + c$$

Hence, for the best fit line, the slope is given as

$$\begin{aligned} m &= \frac{(x_1 - \bar{x})(y_1 - \bar{y}) + (x_2 - \bar{x})(y_2 - \bar{y}) \\ &\quad + (x_3 - \bar{x})(y_3 - \bar{y}) + \dots \\ &\quad + (x_n - \bar{x})(y_n - \bar{y})}{(x_1 - \bar{x})^2 + (x_2 - \bar{x})^2 + (x_3 - \bar{x})^2 + \dots \\ &\quad + (x_n - \bar{x})^2} \\ \Rightarrow m &= \frac{\sum(x_i - \bar{x})(y_i - \bar{y})}{\sum(x_i - \bar{x})^2} \end{aligned}$$

The intercept i.e. c is given by

$$c = \bar{y} - m\bar{x}$$

Note: The differences between a value and the mean value is also referred to as **residuals or errors**.

```
import numpy as np
a1= np.array([1,2,3,4,5])
a2=np.array([10,11,12,13,14])

aa1,aa2,aa11,aa22=train_test_split(a1,a2,test_size=0.2,random_state=42)
aa1

array([5, 3, 1, 4])
```

```
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test= train_test_split(df['enginesize'],df['price'])
```

```
def errors_product():
    prod = (X_train - X_train.mean()) * (y_train - y_train.mean())
    return prod

def squared_errors():
    sq_errors = (X_train - X_train.mean()) ** 2
    return sq_errors

slope = errors_product().sum() / squared_errors().sum()
intercept = y_train.mean() - slope * X_train.mean()

print(f"Slope: {slope} \nIntercept: {intercept}")
```

```
Slope: 165.32203370071696
Intercept: -7590.257181325589
```

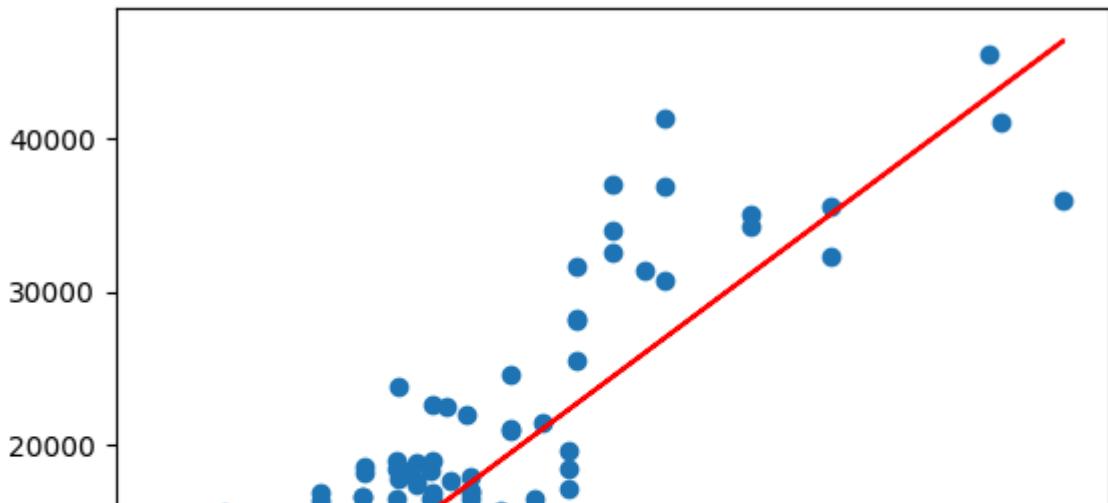
```
X_train_reshaped=X_train.values.reshape(-1,1)
```

```
lr=LinearRegression()
lr.fit(X_train_reshaped,y_train)
lr.intercept_, lr.coef_
```

```
(-7590.257181325582, array([165.3220337]))
```

```
plt.scatter(df['enginesize'],df['price'])
plt.plot(df['enginesize'],slope*df['enginesize']+intercept,color='red')
```

```
[<matplotlib.lines.Line2D at 0x7b69d56bf3a0>]
```



Model Evaluation

❖ The Coefficient of Determination (R-Squared)

The R-squared (R^2) tells us how much of the variance in one variable explains the variance in another variable. It is usually reported in terms of percentage.

Let's compute the coefficient of determination value which is one of the parameters that explains how much variation in one variable can be explained by the other variable through the linear regression model.

$$R^2 = 1 - \frac{\text{SSE}}{\text{SST}}$$

where

$$\text{SSE} = \sum (y - y_{\text{pred}})^2$$

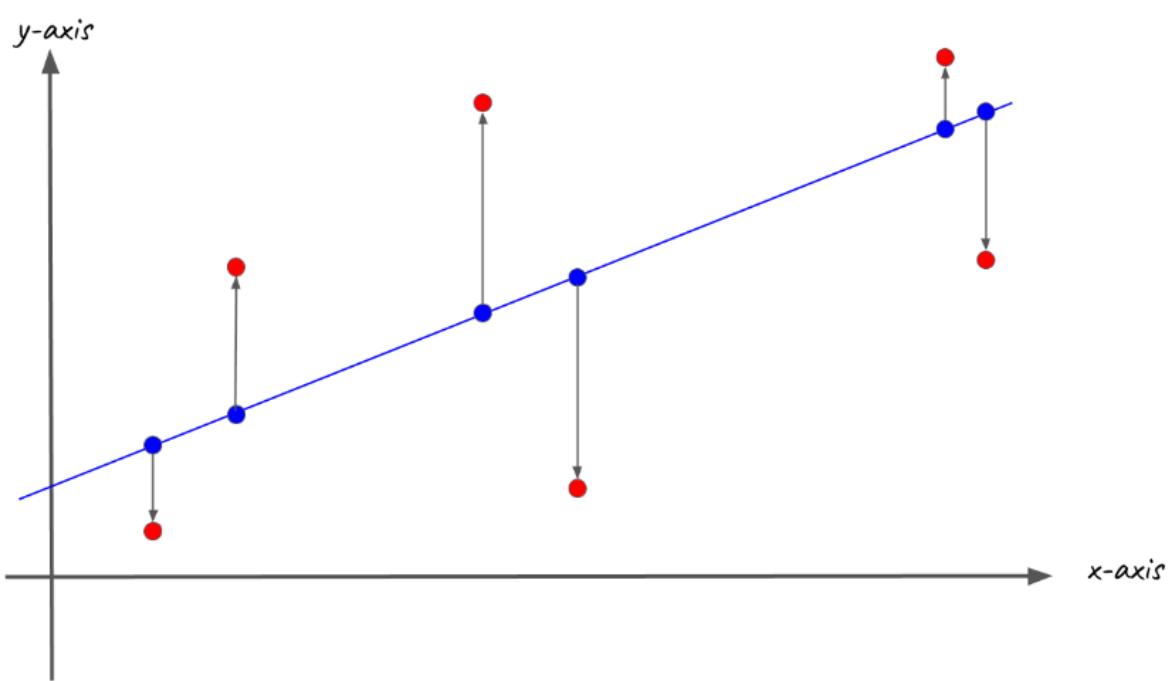
and

$$\text{SST} = \sum (y - \bar{y})^2$$

SSE stands for the sum of squared errors i.e. errors between the actual and the predicted values.

Let there be a straight line which fits them the best

The points marked with the blue colour on the straight line are the corresponding predicted values to the red-coloured points.



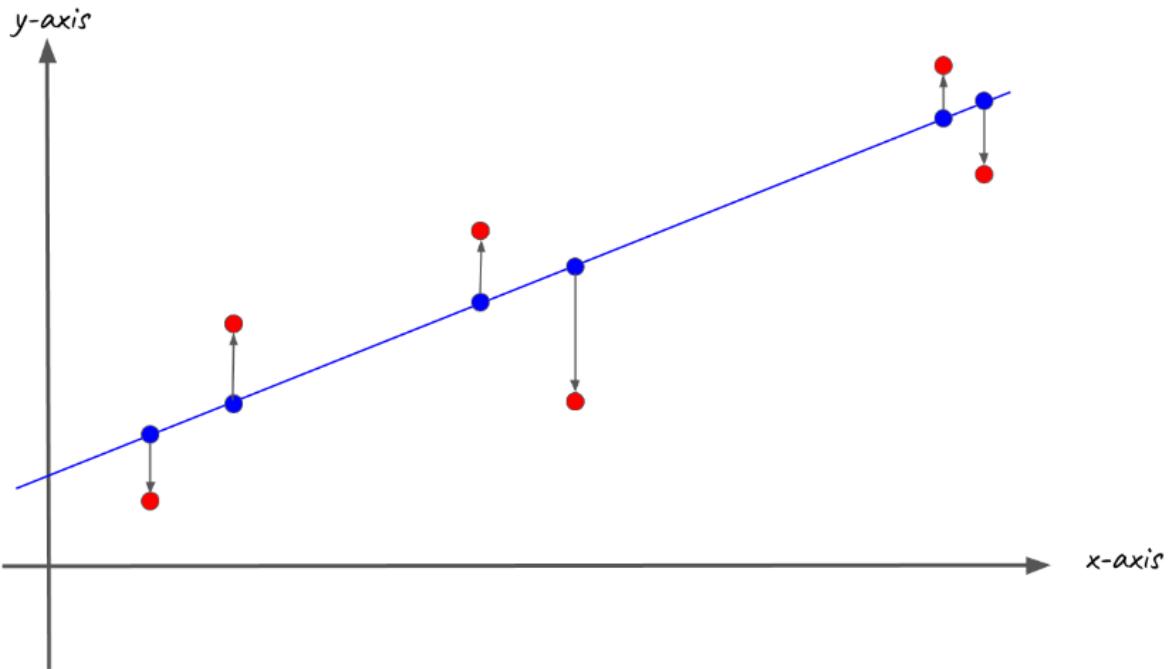
The sum of these distances is the **squared sum of errors (SSE)**. These distances would have been lower if the red-coloured points were more close to the regression line as shown in the image below.

Hence, the SSE value would have been lower. The distance between the actual values and the predicted values are given by

$$|y_i - \hat{y}_i|$$

where

- y_i is the y -coordinate of the actual value and
- \hat{y}_i is the y -coordinate of the corresponding predicted value



Lower the **SSE** value, higher the R^2 value. Higher the R^2 value, better is accuracy.

SST stands for the sum of squared **total** i.e. the errors between the actual values and their mean. Consider the image shown below.

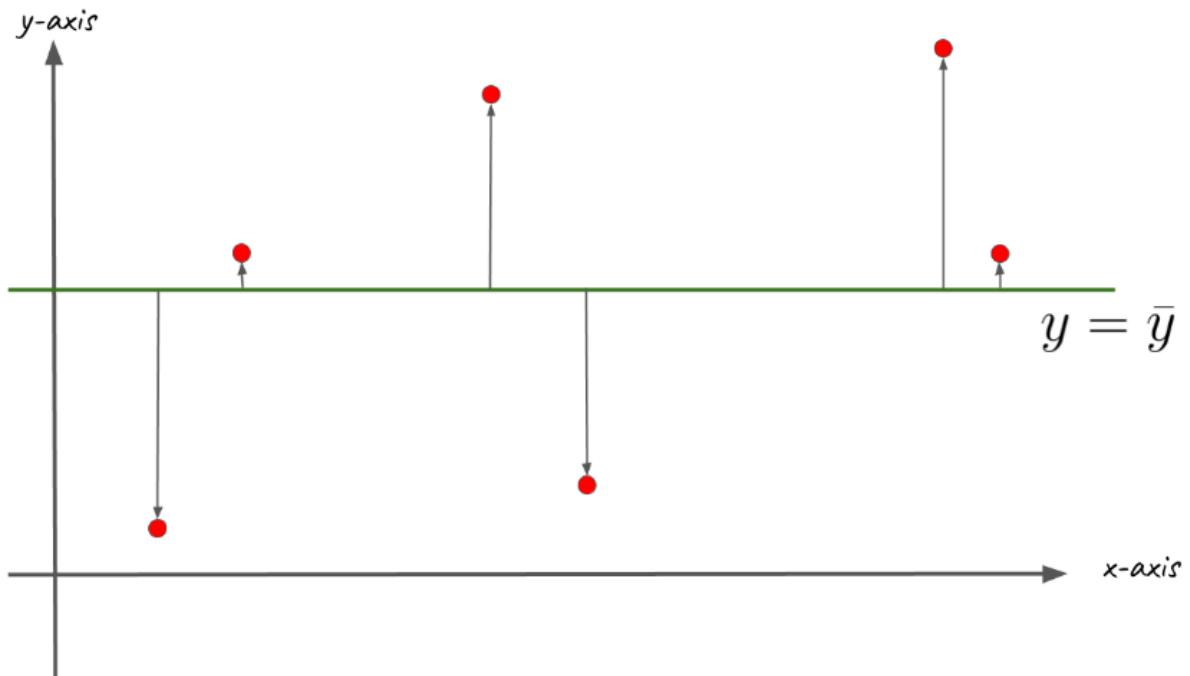
The mean of the actual target variable values i.e. \bar{y} tries to fit all the points. The arrows represent the distances between the points and their mean values.

The sum of these distances is the **squared sum of total (SST)**. They are given by

$$|y_i - \bar{y}|$$

where

- y_i is the y -coordinate of the actual value and
- \bar{y} is their mean value



Also, it is the maximum possible error because the mean line is the worst fit line unless the points follow uniform distribution.

Note:

1. The terms **error**, **residual**, **difference** mean the same thing.
2. It goes without saying that the R^2 value will be between 0 and 1.

```
def r_square(x,y):
    y_pred = slope * x + intercept
    sse = ((y - y_pred) **2).sum()
    sst = ((y - y.mean()) **2).sum()
    r = 1 - (sse/sst)
    return r

print(r_square(X_train, y_train))
print(r_square(X_test,y_test))
```

```
0.7650159366830336
0.7606548315153334
```

```
import numpy as np
np.corrcoef(X_train,y_train)[0,1]**2
```

```
0.7650159366830326
```

```
np.corrcoef(X_train,y_train)
```

```
array([[1.          , 0.87465189],
       [0.87465189, 1.         ]])
```

▼ MSE, RMSE, MAE

Mean Squared Errors (MSE) is the mean of squares of the difference between the actual and the predicted values i.e.

$$\text{MSE} = \frac{1}{n} \sum (y_{\text{actual}} - y_{\text{predicted}})^2$$

where

- y_{actual} is the set of actual values of the target variable
- $y_{\text{predicted}}$ is the set of predicted values of the target variable obtained by deploying some kind of prediction model
- n is the total number of values

Root Mean Squared Errors (RMSE) is the square root of the mean squared errors (MSE) i.e.

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum (y_{\text{actual}} - y_{\text{predicted}})^2}$$

$$\Rightarrow \text{RMSE} = \sqrt{\text{MSE}}$$

Mean Absolute Errors (MAE) is the mean of absolute values of the differences between the actual and the predicted values i.e.

$$\text{MAE} = \frac{1}{n} \sum |y_{\text{actual}} - y_{\text{predicted}}|$$

```
from sklearn.metrics import mean_absolute_error, mean_squared_error, mean_squared_log_error
```

```
X_test_reshaped=X_test.values.reshape(-1,1)
y_train_pred= lr.predict(X_train_reshaped)
y_test_pred= lr.predict(X_test_reshaped)

print("Train")
print(mean_absolute_error(y_train,y_train_pred))
print(mean_squared_error(y_train,y_train_pred))
print(mean_squared_log_error(y_train,y_train_pred))
print()
print('Test')

print(mean_absolute_error(y_test,y_test_pred))
print(mean_squared_error(y_test,y_test_pred))
print(mean_squared_log_error(y_test,y_test_pred))
```

```
Train
2906.3830116569407
14675971.645771094
0.07982374537928014
```

```
Test
2773.9908161019634
15661604.54844862
0.09333552536979671
```

✓ Residual Analysis

In the residual analysis, you need to check if the error terms are normally distributed (which is infact, one of the major assumptions of linear regression). Why? Because, formally, a simple linear regression model is given as

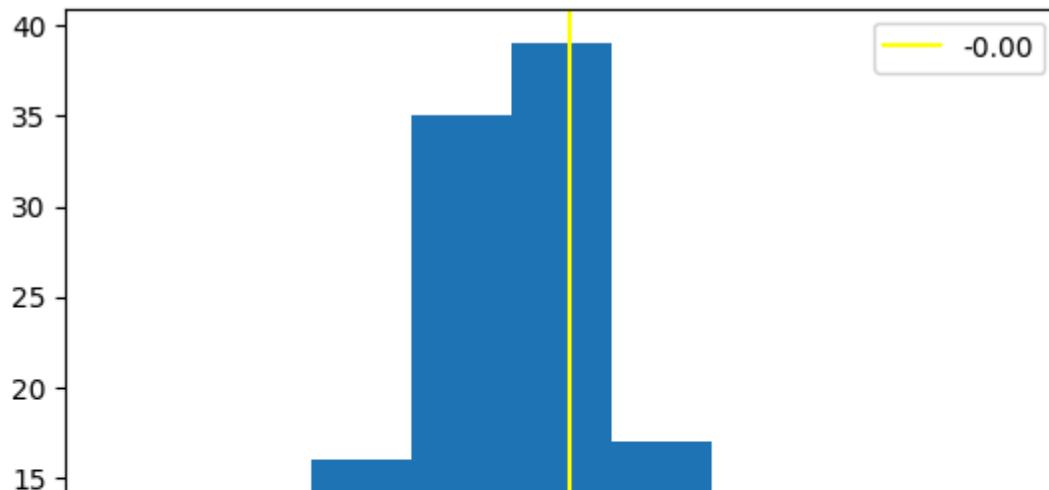
$$Y = \beta_0 + \beta_1 x + \epsilon$$

where

- x is the independent variable
- Y is the response to the independent variable (or predicted value or dependent variable)
- β_0 (intercept made by the best fit line with the y -axis) and β_1 (slope of the best fit line) are called regression coefficients
- ϵ is the random error obtained along with the predicted value

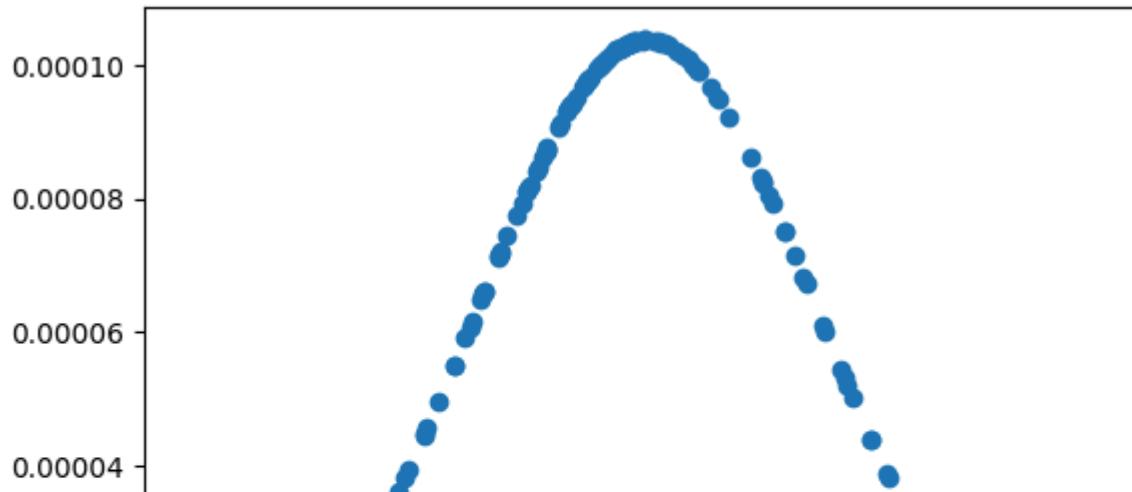
For a line to be the best fit line, the mean of random errors i.e. mean of ϵ should be 0.

```
# Residual Analysis
train_error= y_train- y_train_pred
plt.hist(train_error, bins='sturges')
plt.axvline(train_error.mean(),label=f'{train_error.mean():.2f}',color='yellow')
plt.legend()
plt.show()
```

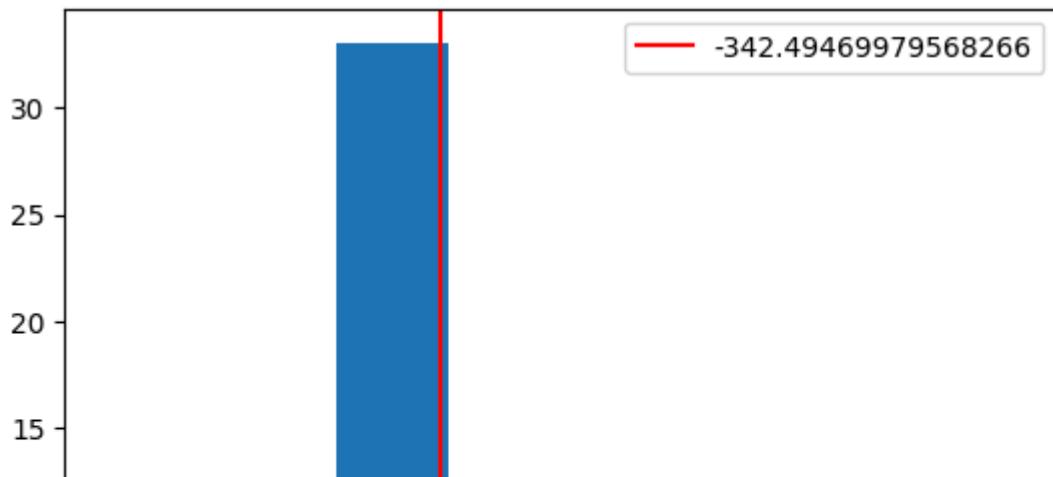


```
from scipy.stats import norm
d = norm.pdf(train_error,train_error.mean(),train_error.std())
plt.scatter(train_error,d)
```

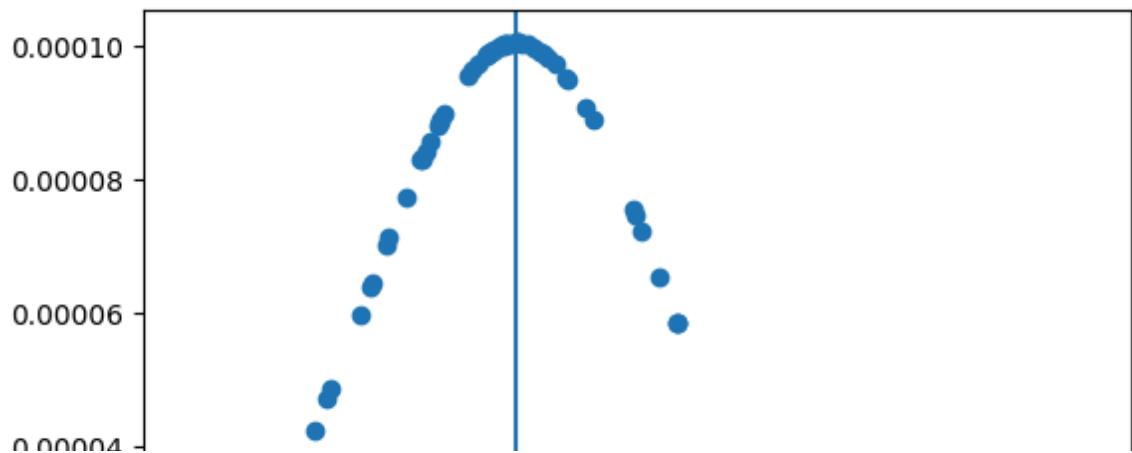
```
<matplotlib.collections.PathCollection at 0x7b69d5601c60>
```



```
test_error= y_test-y_test_pred
plt.hist(test_error,bins='sturges')
plt.axvline(test_error.mean(),label= f'{test_error.mean()}',color='red')
plt.legend()
plt.show()
```



```
d2= norm.pdf(test_error,test_error.mean(),test_error.std())
plt.scatter(test_error,d2)
plt.axvline(test_error.mean())
plt.show()
```

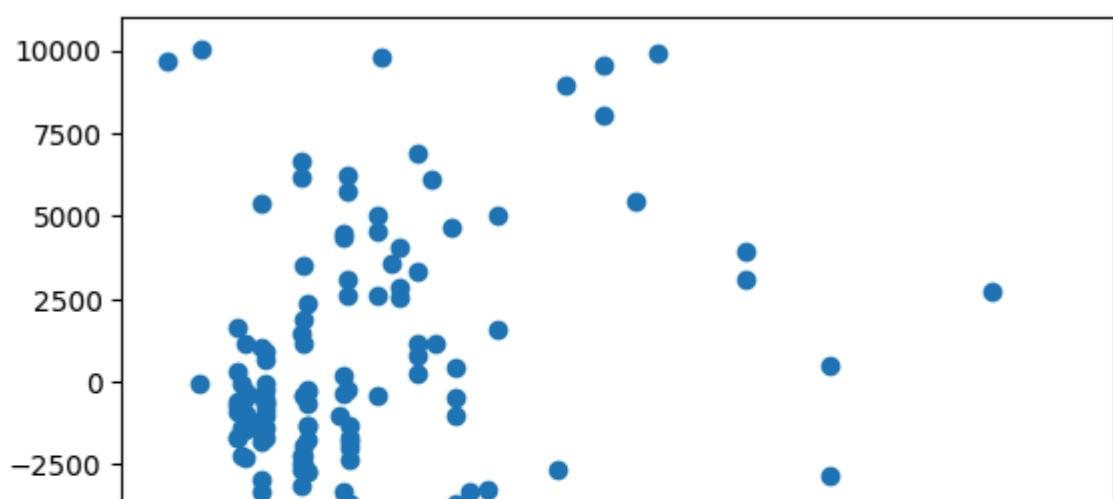


Homoscedasticity and Heteroscedasticity

Check for the trend in the scatter plot between the errors and the feature and target variables. There should not be a trend.

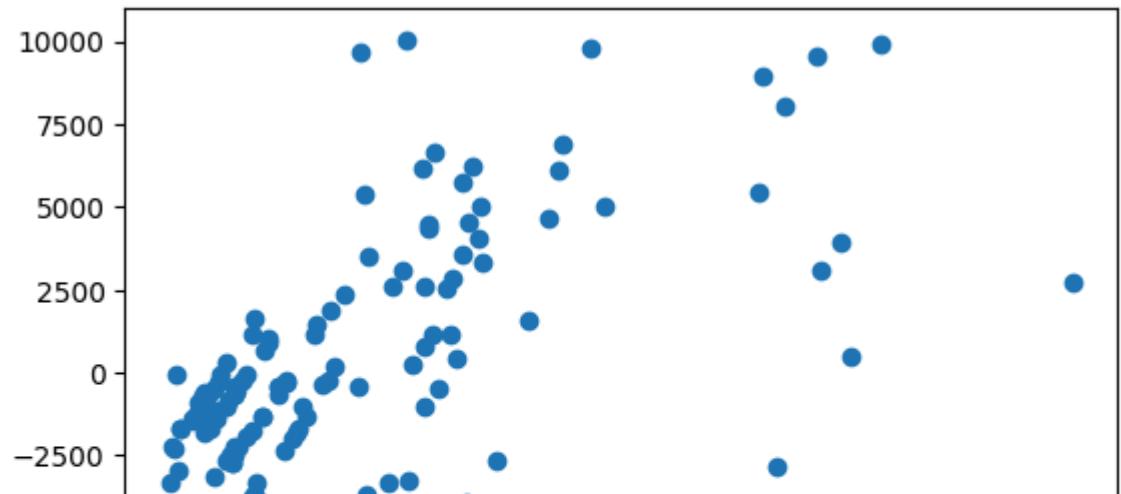
```
plt.scatter(X_train,train_error)
```

<matplotlib.collections.PathCollection at 0x7b69d5731b70>



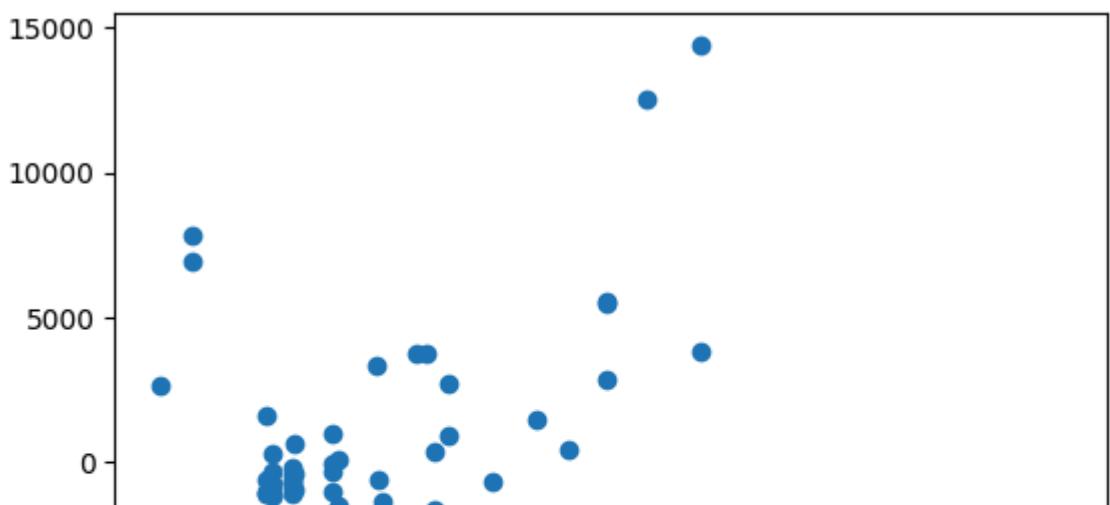
```
plt.scatter(y_train,train_error)
```

```
<matplotlib.collections.PathCollection at 0x7b69d54e2b30>
```



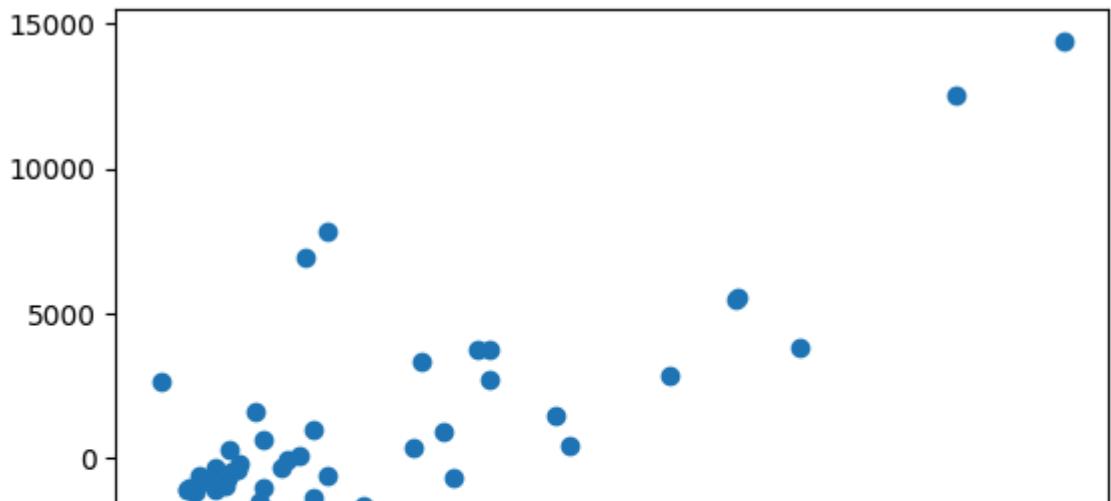
```
plt.scatter(X_test,test_error)
```

```
<matplotlib.collections.PathCollection at 0x7b69d557f100>
```



```
plt.scatter(y_test,test_error)
```

```
<matplotlib.collections.PathCollection at 0x7b69d53eb0>
```



▼ Multiple Linear Regression

▼ Preparing data

Data Scaling

```
from sklearn.preprocessing import StandardScaler
X= df_new.drop('price',axis=1)
y= df_new['price']
X_train,X_test,y_train,y_test= train_test_split(X,y,test_size=0.3,random_state=42)
X_train.head()
```

	symboling	doornumber	wheelbase	carlength	carwidth	carheight	curbwe
177	-1	4	102.4	175.6	66.5	53.9	1750
75	1	2	102.7	178.4	68.0	54.8	1750
174	-1	4	102.4	175.6	66.5	54.9	1750
31	2	2	86.6	144.6	63.9	50.8	1750
12	0	2	101.2	176.8	64.8	54.3	1750

5 rows × 59 columns

X.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 59 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   symboling        205 non-null    int64  
 1   doornumber       205 non-null    int64  
 2   wheelbase        205 non-null    float64 
 3   carlength        205 non-null    float64 
 4   carwidth         205 non-null    float64 
 5   carheight        205 non-null    float64 
 6   curbweight       205 non-null    int64  
 7   cylindernumber   205 non-null    int64  
 8   enginesize       205 non-null    int64  
 9   boreratio        205 non-null    float64 
 10  stroke          205 non-null    float64 
 11  compressionratio 205 non-null    float64 
 12  horsepower       205 non-null    int64  
 13  peakrpm          205 non-null    int64  
 14  citympg          205 non-null    int64  
 15  highwaympg       205 non-null    int64  
 16  fueltype_gas     205 non-null    int64  
 17  aspiration_turbo 205 non-null    int64  
 18  carbody_hardtop  205 non-null    int64  
 19  carbody_hatchback 205 non-null    int64  
 20  carbody_sedan    205 non-null    int64  
 21  carbody_wagon    205 non-null    int64  
 22  drivewheel_fwd   205 non-null    int64  
 23  drivewheel_rwd   205 non-null    int64  
 24  enginelocation_rear 205 non-null    int64  
 25  enginetype_dohcv 205 non-null    int64  
 26  enginetype_l      205 non-null    int64  
 27  enginetype_ohc    205 non-null    int64  
 28  enginetype_ohcf   205 non-null    int64  
 29  enginetype_ocv    205 non-null    int64  
 30  enginetype_rotor  205 non-null    int64  
 31  fuelsystem_2bbl   205 non-null    int64  
 32  fuelsystem_4bbl   205 non-null    int64  
 33  fuelsystem_idi    205 non-null    int64  
 34  fuelsystem_mfi    205 non-null    int64  
 35  fuelsystem_mpfi   205 non-null    int64  
 36  fuelsystem_spdi   205 non-null    int64
```

37	fuelsystem_spfi	205	non-null	int64
38	CarCompany_audi	205	non-null	int64
39	CarCompany_bmw	205	non-null	int64
40	CarCompany_buick	205	non-null	int64
41	CarCompany_chevrolet	205	non-null	int64
42	CarCompany_dodge	205	non-null	int64
43	CarCompany_honda	205	non-null	int64
44	CarCompany_isuzu	205	non-null	int64
45	CarCompany_jaguar	205	non-null	int64
46	CarCompany_mazda	205	non-null	int64
47	CarCompany_mercury	205	non-null	int64
48	CarCompany_mitsubishi	205	non-null	int64
49	CarCompany_nissan	205	non-null	int64
50	CarCompany_peugeot	205	non-null	int64
51	CarCompany_plymouth	205	non-null	int64
52	CarCompany_porsche	205	non-null	int64

```
X_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 143 entries, 177 to 102
Data columns (total 59 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   symboling        143 non-null    int64  
 1   doornumber       143 non-null    int64  
 2   wheelbase        143 non-null    float64 
 3   carlength        143 non-null    float64 
 4   carwidth         143 non-null    float64 
 5   carheight        143 non-null    float64 
 6   curbweight       143 non-null    int64  
 7   cylindernumber   143 non-null    int64  
 8   enginesize       143 non-null    int64  
 9   boreratio        143 non-null    float64 
 10  stroke          143 non-null    float64 
 11  compressionratio 143 non-null    float64 
 12  horsepower       143 non-null    int64  
 13  peakrpm          143 non-null    int64  
 14  citympg          143 non-null    int64  
 15  highwaympg       143 non-null    int64  
 16  fueltype_gas     143 non-null    int64  
 17  aspiration_turbo 143 non-null    int64  
 18  carbody_hardtop  143 non-null    int64  
 19  carbody_hatchback 143 non-null    int64  
 20  carbody_sedan    143 non-null    int64  
 21  carbody_wagon    143 non-null    int64  
 22  drivewheel_fwd   143 non-null    int64  
 23  drivewheel_rwd   143 non-null    int64  
 24  enginelocation_rear 143 non-null    int64  
 25  enginetype_dohcv  143 non-null    int64  
 26  enginetype_l      143 non-null    int64  
 27  enginetype_ohc    143 non-null    int64  
 28  enginetype_ohcf   143 non-null    int64  
 29  enginetype_ohcv   143 non-null    int64  
 30  enginetype_rotor  143 non-null    int64  
 31  fuelsystem_2bbl   143 non-null    int64  
 32  fuelsystem_4bbl   143 non-null    int64  
 33  fuelsystem_idi    143 non-null    int64  
 34  fuelsystem_mfi    143 non-null    int64  
 35  fuelsystem_mpfi   143 non-null    int64
```

36	fuelsystem_spdi	143	non-null	int64
37	fuelsystem_spfi	143	non-null	int64
38	CarCompany_audi	143	non-null	int64
39	CarCompany_bmw	143	non-null	int64
40	CarCompany_buick	143	non-null	int64
41	CarCompany_chevrolet	143	non-null	int64
42	CarCompany_dodge	143	non-null	int64
43	CarCompany_honda	143	non-null	int64
44	CarCompany_isuzu	143	non-null	int64
45	CarCompany_jaguar	143	non-null	int64
46	CarCompany_mazda	143	non-null	int64
47	CarCompany_mercury	143	non-null	int64
48	CarCompany_mitsubishi	143	non-null	int64
49	CarCompany_nissan	143	non-null	int64
50	CarCompany_peugeot	143	non-null	int64
51	CarCompany_plymouth	143	non-null	int64
52	CarCompany_porsche	143	non-null	int64

```
ss= StandardScaler()
ss_values= ss.fit_transform(X_train[X_train.columns[:16]])
X_train[X_train.columns[:16]]= ss_values
X_train.head()
```

	symboling	doornumber	wheelbase	carlength	carwidth	carheight	curbwe
177	-1.5000	0.887412	0.573309	0.076413	0.235105	0.043859	-0.22
75	0.1250	-1.126872	0.622875	0.302880	0.924984	0.408026	0.64
174	-1.5000	0.887412	0.573309	0.076413	0.235105	0.448489	-0.17
31	0.9375	-1.126872	-2.037199	-2.430901	-0.960684	-1.210497	-1.44
12	-0.6875	-1.126872	0.375042	0.173470	-0.546757	0.205711	0.26

5 rows × 59 columns

```
import statsmodels.api as sm
X_train_sm= sm.add_constant(X_train)
model= sm.OLS(y_train,X_train_sm).fit()
model.params
```

	θ
const	1.111326e+04
symboling	-2.583520e+02
doornumber	2.554028e+02
wheelbase	4.605784e+02
carlength	-6.252099e+02
carwidth	1.483336e+03
carheight	-3.289250e+02
curbweight	2.130519e+03
cylindernumber	-8.448995e+02
enginesize	4.068721e+03
boreratio	-1.024074e+03
stroke	-7.567151e+01
compressionratio	4.636386e+02
horsepower	2.553877e+02
peakrpm	8.288722e+02
citympg	6.368825e+02
highwaympg	-2.391232e+02
fueltype_gas	6.290680e+03
aspiration_turbo	2.528179e+03
carbody_hardtop	-2.337527e+02
carbody_hatchback	-3.517083e+03
carbody_sedan	-3.523827e+03
carbody_wagon	-4.154162e+03
drivewheel_fwd	-8.505823e+02
drivewheel_rwd	-1.393007e+03
enginelocation_rear	4.920851e+03
enginetype_dohcv	2.908615e+02
enginetype_l	-1.259194e+03
enginetype_ohc	-3.180058e+02
enginetype_ohcf	1.138945e+03
enainetvne ohcv	-1.391856e+03

```
print(model.summary())
```

fuelsystem_2bbl	2.404488e+03	OLS Regression Results				
Dep. Variable:	fuelsystem_4bbl	1.132048e+03	price	R-squared:	0.9	
Model:	fuelsystem_idi	4.822583e+03	OLS	Adj. R-squared:	0.9	
Method:		Least Squares		F-statistic:	67	
Date:	fuelsystem_mfi	Sat, 07.02.2023 00:24:42		Prob (F-statistic):	3.52e-12	
Time:		03:28:27		Log-Likelihood:	-1214	
No. Observations:	fuelsystem_mpfi	1.568049e+03		AIC:	254	
Df Residuals:		87		BIC:	276	
Df Model:	fuelsystem_spdi	1.158155e+03				
Covariance Type:	fuelsystem_spfi	nonrobust				
		2.039415e+03				
CarCompany_audi	1.496863e+01	coef	std err	t	P> t	[0.0]
const		1.11272259e+03	2193.733	5.066	0.000	6752.98
CarCompany_bmw		-258.3520	337.076	-0.766	0.445	-928.31
CarCompany_buick		256.875953e+03	235.521	0.927	0.357	-292.21
wheelbase		460.5784	667.071	0.690	0.492	-865.29
CarCompany_chevrolet		-625.2099	673.875	-0.928	0.356	-1964.66
carlength		1483.3360	501.492	2.958	0.004	486.56
CarCompany_dodge		-328.9250	360.885	-0.911	0.365	-1046.21
curbweight		2130.5189	867.221	2.457	0.016	406.81
CarCompany_honda		-844.8995	720.653	-1.172	0.244	-2277.21
cylindernumber		4068.521195e+03	1079.175	3.770	0.000	1923.74
CarCompany_isuzu		-1024.0736	376.775	-2.718	0.008	-1772.91
boreratio		-75.491049e+03	232.058	-0.268	0.789	-636.29
CarCompany_jaguar		463.6386	1675.384	0.277	0.783	-2866.36
compressionratio		255.587720e+03	829.610	0.308	0.759	-1393.51
horsepower		828.8722	325.351	2.548	0.013	182.26
peakrpm		636.8825	850.582	0.749	0.456	-1053.74
citympg		1.594039e+03				
CarCompany_mercury		-239.1232	766.429	-0.312	0.756	-1762.48
CarCompany_mitsubishi		6290.6800	2516.294	2.500	0.014	1289.21
fueltypes_gas		2529.072283e+03	820.112	3.083	0.003	898.11
aspire		-233.7527	1674.144	-0.140	0.889	-3561.29
CarCompany_nissan		-3517.283194e+03	136.021	-3.096	0.003	-5775.04
carbody_hardtop		-3523.8275	1243.949	-2.833	0.006	-5996.31
CarCompany_peugeot		4154.416191e+03	1363.304	-3.047	0.003	-6863.81
carbody_sedan		-850.5823	1150.963	-0.739	0.462	-3138.24
CarCompany_plymouth		-1393.0072	1327.046	-1.050	0.297	-4030.61
drivewheel_fwd		4920.8510	1878.236	2.620	0.010	1187.61
CarCompany_porsche		290.8615	3118.387	0.093	0.926	-5907.21
drivewheel_rwd		3.693763e+03				
engine_location_rear		-1258.194085e+02	958.461	-1.314	0.192	-3164.24
CarCompany_renault		-318.0058	902.080	-0.353	0.725	-2110.98
enginetype_dohcv		1138.7941206e+03	1059.824	1.075	0.286	-967.51
enginetype_l		-1391.8556	1301.970	-1.069	0.288	-3979.66
CarCompany_subaru		6742.7581770e+03	201.565	2.707	0.008	1792.13
enginetype_ohcv		2404.4879	1359.986	1.768	0.081	-298.61
CarCompany_toyota		1132.0482	2592.199	0.437	0.663	-4020.21
fuelsystem_2bbl		4822.5828	4052.732	1.190	0.237	-3232.66
CarCompany_volkswagen		-1.027e-12	1.89e-12	-0.543	0.589	-4.79e-1
fuelsystem_4bbl		1568.0491	1482.816	1.057	0.293	-1379.26
CarCompany_idi		1158.1552	1718.412	0.674	0.502	-2257.31
fuelsystem_mpfi		2039.4149	2480.546	0.822	0.413	-2890.91
CarCompany_mfi		14.9686	2053.835	0.007	0.994	-4067.21
CarCompany_bmw		7272.2594	1965.043	3.701	0.000	3366.51
CarCompany_buick		6875.9526	2219.913	3.097	0.003	2463.61
CarCompany_chevrolet		-4172.9890	2343.759	-1.780	0.078	-8831.46
		1110.1006	1700.100	0.200	0.800	-7000.00

Adjusted R^2

In the case of multiple linear regression, **adjusted R^2** value takes precedence over the R^2 value. It is calculated as:

$$R_{\text{adj}}^2 = 1 - \frac{(1 - R^2)(N - 1)}{N - p - 1}$$

where

- R^2 is the coefficient of determination
- N is number of instances (or rows) in the dataset
- p is the number of independent variables (excluding constant) in the dataset

the R_{adj}^2 will always be less than or equal to the R^2 value i.e. $R_{\text{adj}}^2 \leq R^2$.

Why adjusted R-squared is a better metric in multiple linear regression?

As you add more and more independent variables, the R^2 squared values increases even if the independent variable has no contribution in predicting the values of the target variable. Hence, the adjusted R^2 value penalises the unnecessary inclusion of more independent variables.

So, if adding more independent (or feature) variables leads to an increase in the adjusted R^2 value, then it is a good sign. However, if adding more independent (or feature) variables leads to a decrease in the adjusted R^2 value, it is a bad sign.

In this case, the R_{adj}^2 is quite high but the p-values for many of the columns is greater than 0.05 which is not a good sign. It means, these variables are insignificant in predicting the price of a car. Also, if we calculate variance inflation factor values for these columns, they would be very very high than 10.

Ordinary Least Squares (OLS)

Consider the regression equation

$$Y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \dots + \beta_k x_k + \epsilon$$

where

- $x_1, x_2, x_3, \dots, x_k$ are independent variables or features
- Y is the response to the independent variable (or predicted value or dependent variable)
- $\beta_0, \beta_1, \beta_2, \dots, \beta_k$ are the corresponding regression coefficients of the independent variables

- ϵ is the random error obtained along with the predicted value which follows normal distribution with mean 0 and some standard deviation of σ

The parameters $\beta_0, \beta_1, \beta_2, \dots, \beta_n$ and σ are assumed to be unknown and must be estimated from the data, which we shall suppose will consist of the values of $Y_1, Y_2, Y_3, \dots, Y_n$ where Y_i is the response level corresponding to the k features $x_{i1}, \dots, x_{i2}, \dots, x_{ik}$. That is, the Y_i are related to these features through

$$\begin{aligned} E[Y_i] &= \beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \beta_3 x_{i3} + \dots \\ &\quad + \beta_k x_{ik} \end{aligned}$$

where

- $E[Y_i]$ means **expected value** for an instance i . In simple terms, instance or i denotes a row in a data frame
- x_{i1} denotes item at the i^{th} row in the 1^{st} column in a data frame having only features
- x_{i2} denotes item at the i^{th} row in the 2^{nd} column in a data frame having only features
- x_{i3} denotes item at the i^{th} row in the 3^{rd} column in a data frame having only features
- \dots
- x_{ik} denotes item at the i^{th} row in the k^{th} column in a data frame having only features

As we said earlier, the difference between the actual and the predicted values should be 0 or close to 0 for an accurate prediction model i.e.

$$Y_1 - E[Y_1] \approx 0$$

$$Y_2 - E[Y_2] \approx 0$$

$$Y_3 - E[Y_3] \approx 0$$

⋮

$$Y_N - E[Y_N] \approx 0$$

where N is the total number of instances (or rows in a data frame).

The OLS says that the sum of squares of all these errors i.e.

$$\begin{aligned} J &= (Y_1 - E[Y_1])^2 + (Y_2 - E[Y_2])^2 \\ &\quad + (Y_3 - E[Y_3])^2 + \dots + (Y_N - E[Y_N])^2 \end{aligned}$$

should be the least or minimum.

The above expression can be compressed as

$$J = \sum_{i=1}^N (Y_i - E[Y_i])^2$$

So in general, it can be written as

$$\begin{aligned} J(\beta, x) = & \sum_{i=1}^N (Y_i - \beta_0 - \beta_1 x_{i1} - \beta_2 x_{i2} \\ & - \beta_3 x_{i3} - \cdots - \beta_k x_{ik})^2 \end{aligned}$$

where $J(\beta, x)$ denotes the sum of the squared errors is dependent on the coefficients $(\beta_0, \beta_1, \beta_2, \dots, \beta_k)$ and features $(x_1, x_2, x_3, \dots, x_k)$

To find the points of maxima (peak) or minima (valley), we differentiate a mathematical function w.r.t. independent variable and equate the result obtained to 0 because the slope of a curve at the point of maxima (peak) or minima (valley) is 0. Differentiation (or derivative) represents slope at a point.

In the above equation, all the x quantities are known quantities as we have seen earlier. So the $\beta_0, \beta_1, \beta_2, \dots, \beta_k$ are unknown quantities. Thus, they are independent variables.

Here $\beta_0, \beta_1, \beta_2, \dots, \beta_k$ are independent of each other. Hence, we can do partial differentiation w.r.t. to each of the betas independently.

Let's differentiate $J(\beta, x)$ w.r.t. β_0 . So every other term apart from β_0 will be treated as a constant. And the differentiation (or derivative) of a constant is 0.

$$\begin{aligned} \frac{\partial J}{\partial \beta_0} &= 2 \sum_{i=1}^N (Y_i - \beta_0 - \beta_1 x_{i1} - \beta_2 x_{i2} \\ &\quad - \beta_3 x_{i3} - \cdots - \beta_k x_{ik})(-1) = 0 \\ \Rightarrow & \sum_{i=1}^N (Y_i - \beta_0 - \beta_1 x_{i1} - \beta_2 x_{i2} - \beta_3 x_{i3} - \\ &\quad \cdots - \beta_k x_{ik}) = 0 \end{aligned}$$

Similarly,

$$\begin{aligned} \frac{\partial J}{\partial \beta_1} &= 2 \sum_{i=1}^N (Y_i - \beta_0 - \beta_1 x_{i1} - \beta_2 x_{i2} \\ &\quad - \beta_3 x_{i3} - \cdots - \beta_k x_{ik})(-x_{i1}) = 0 \\ \Rightarrow & \sum_{i=1}^N (Y_i - \beta_0 - \beta_1 x_{i1} - \beta_2 x_{i2} - \beta_3 x_{i3} - \\ &\quad \cdots - \beta_k x_{ik})x_{i1} = 0 \end{aligned}$$

$$\begin{aligned}
 \frac{\partial J}{\partial \beta_2} &= 2 \sum_{i=1}^N (Y_i - \beta_0 - \beta_1 x_{i1} - \beta_2 x_{i2} \\
 &\quad - \beta_3 x_{i3} - \cdots - \beta_k x_{ik})(-x_{i2}) = 0 \\
 \Rightarrow \sum_{i=1}^N (Y_i - \beta_0 - \beta_1 x_{i1} - \beta_2 x_{i2} - \beta_3 x_{i3} - \\
 &\quad \cdots - \beta_k x_{ik}) x_{i2} = 0 \\
 &\quad \vdots \\
 \frac{\partial J}{\partial \beta_k} &= 2 \sum_{i=1}^N (Y_i - \beta_0 - \beta_1 x_{i1} - \beta_2 x_{i2} \\
 &\quad - \beta_3 x_{i3} - \cdots - \beta_k x_{ik})(-x_{ik}) = 0 \\
 \Rightarrow \sum_{i=1}^N (Y_i - \beta_0 - \beta_1 x_{i1} - \beta_2 x_{i2} - \beta_3 x_{i3} - \\
 &\quad \cdots - \beta_k x_{ik}) x_{ik} = 0
 \end{aligned}$$

On further reducing the above $k + 1$ equations, we get

$$\begin{aligned}
 \sum_{i=1}^N Y_i &= N\beta_0 + \beta_1 \sum_{i=1}^N x_{i1} + \beta_2 \sum_{i=1}^N x_{i2} + \\
 &\quad \cdots + \beta_k \sum_{i=1}^N x_{ik} \\
 \sum_{i=1}^N Y_i x_{i1} &= \beta_0 \sum_{i=1}^N x_{i1} + \beta_1 \sum_{i=1}^N x_{i1}^2 + \beta_2 \\
 &\quad \sum_{i=1}^N x_{i1} x_{i2} + \cdots + \beta_k \sum_{i=1}^N x_{i1} x_{ik} \\
 &\quad \vdots \\
 \sum_{i=1}^N Y_i x_{ik} &= \beta_0 \sum_{i=1}^N x_{ik} + \beta_1 \sum_{i=1}^N x_{ik} x_{i1} \\
 &\quad + \beta_2 \sum_{i=1}^N x_{ik} x_{i2} + \cdots + \beta_k \sum_{i=1}^N x_{ik}^2
 \end{aligned}$$

Now we have $k + 1$ linear equations having $k + 1$ unknowns i.e. $\beta_0, \beta_1, \beta_2, \dots, \beta_k$. By solving these $k + 1$ equations, we can get the beta values. This is exactly the same as solving two linear equations having two unknowns. For e.g., the solution to the two linear equations

$$8\beta_0 + 7\beta_1 = 38 \text{ and } 3\beta_0 - 5\beta_1 = -1$$

is

$$\beta_0 = 3 \text{ and } \beta_1 = 2$$

So all-in-all, **ordinary least squares** says that **find the values of the coefficients ($\beta_0, \beta_1, \beta_2, \dots, \beta_k$) such that the sum of the squares of differences between the actual values and the predicted values is minimum.**

To solve $k + 1$ linear equations having $k + 1$ unknowns, you need to know matrices.

The above $k + 1$ linear equations can also be written as

$$\begin{aligned} & \$\$ \begin{bmatrix} \sum_{i=1}^N Y_i \\ \sum_{i=1}^N Y_i x_{i1} \\ \sum_{i=1}^N Y_i x_{i2} \\ \vdots \\ \sum_{i=1}^N Y_i x_{ik} \end{bmatrix} = \begin{bmatrix} \sum_{i=1}^N 1 & \sum_{i=1}^N x_{i1} & \dots & \sum_{i=1}^N x_{ik} \\ \sum_{i=1}^N x_{i1} & \sum_{i=1}^N x_{i1}^2 & \dots & \sum_{i=1}^N x_{i1} x_{i2} \\ \vdots & \vdots & \ddots & \vdots \\ \sum_{i=1}^N x_{ik} & \sum_{i=1}^N x_{i2} & \dots & \sum_{i=1}^N x_{i1} x_{i2} \\ \vdots & \vdots & \ddots & \vdots \\ \sum_{i=1}^N x_{ik} & \sum_{i=1}^N x_{ik} x_{i1} & \dots & \sum_{i=1}^N x_{ik} x_{ik} \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \\ \vdots \\ \beta_k \end{bmatrix} \\ & \end{bmatrix} \\ & \end{aligned}$$

in the matrix form. The above matrix equation can also be written as

$$X^T Y = X^T X B$$

or

$$X^T X B = X^T Y$$

where

$$\begin{aligned} X = & \begin{bmatrix} 1 & x_{11} & x_{12} & x_{13} & \dots & x_{1k} \\ 1 & x_{21} & x_{22} & x_{23} & \dots & x_{2k} \\ 1 & x_{31} & x_{32} & x_{33} & \dots & x_{3k} \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_{N1} & x_{N2} & x_{N3} & \dots & x_{Nk} \end{bmatrix} \\ & \end{aligned}$$

$$\begin{aligned} X^T = & \begin{bmatrix} 1 & 1 & 1 & \dots & 1 \\ x_{11} & x_{21} & x_{31} & \dots & x_{N1} \\ x_{12} & x_{22} & x_{32} & \dots & x_{N2} \\ x_{13} & x_{23} & x_{33} & \dots & x_{N3} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ x_{1k} & x_{2k} & x_{3k} & \dots & x_{Nk} \end{bmatrix} \\ & \end{aligned}$$

$$\begin{aligned} Y = & \begin{bmatrix} Y_1 \\ Y_2 \\ Y_3 \\ \vdots \\ Y_N \end{bmatrix} \\ B = & \begin{bmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \\ \vdots \\ \beta_k \end{bmatrix} \\ & \end{bmatrix} \\ & \end{aligned}$$

In the matrix equation,

$$X^T X B = X^T Y$$

to obtain only the matrix B on the left-hand side, you need to multiply both the sides by $(X^T X)^{-1}$, i.e.

$$\$(X^T X)^{-1} X^T X B = (X^T X)^{-1} X^T Y\$\$$$

To simplify the above equation, let

$$\$Z = X^T X\$$$

$$\$\therefore Z^{-1} = (X^T X)^{-1}\$\$$$

Hence, the above equation becomes

$$\$Z^{-1} Z B = Z^{-1} X^T Y\$\$$$

$$\$\Rightarrow IB = Z^{-1} X^T Y \quad (\text{because } Z^{-1} Z = I)\$\$$$

$$\$\Rightarrow B = Z^{-1} X^T Y \quad (\text{because } IB = B)\$\$$$

$$\text{Let } U = X^T Y\$\$$$

$$\$\therefore B = Z^{-1} U\$\$$$

Now, you need to obtain the Z^{-1} and multiply it with the matrix U to estimate the values of betas using the matrix operations only. But before that, you need to add a new column to the matrix X , i.e., `X_train`. All the items of this new column should be \$1\$.

Variance Inflation factor (VIF)

Measure of Multicollinearity

Variance Inflation Factor (VIF) is a way to detect multicollinearity between independent variables in a dataset. We calculate the VIF values to measure the extent of multicollinearity between the independent variables.

For k different independent variables, we can calculate k different VIFs (one for each x_i where $i = 1, 2, 3, \dots, k$) in three steps:

Step one

First, build a multiple linear regression model wherein x_i is a target variable and it is a function of all the other feature variables as illustrated in the equation below.

$$\$x_1 = \beta_0 + \beta_1 x_2 + \beta_2 x_3 + \beta_3 x_4 + \dots + \beta_k x_k + \epsilon\$\$$$

Here,

- x_1 is a feature acting as the target (or dependent) variable in above equation
- $x_2, x_3, x_4, \dots, x_k$ are independent variables or features
- $\beta_0^*, \beta_2^*, \beta_3^*, \dots, \beta_k^*$ are the corresponding regression coefficients of the independent variables in the above linear regression equation
- ϵ^* is the random error obtained along with the predicted value

Step two

Then, calculate the VIF for x_i using the following formula:

$$\text{VIF}_i = \frac{1-R_i^2}{1-R_{-i}^2}$$

where R^2_i is the coefficient of determination of the regression equation in step one, with x_i on the left hand side, and all other independent variables on the right hand side.

Step three

Analyse the extent of multicollinearity by considering the magnitude of the VIF_i . A rule of thumb is that if $\text{VIF}_i > 10$, then multicollinearity is high. In that case, the x_i feature must be dropped to predict the values of the target (or dependent) variable. A cutoff of 5 is also commonly used.

```
from statsmodels.stats.outliers_influence import variance_inflation_factor
X_train_new= X_train[['enginesize','curbweight','horsepower','citympg']]
X_train_new_sm= sm.add_constant(X_train_new)
model_new= sm.OLS(y_train,X_train_new_sm)
vif_df= pd.DataFrame()
vif_df['Features']= X_train_new_sm.columns
vif_df['VIF']= [variance_inflation_factor(X_train_new_sm.values, i) for i in
vif_df]
```

	Features	VIF
0	const	1.000000
1	enginesize	4.901158
2	curbweight	4.732600
3	horsepower	4.533126
4	citympg	3.942814

Calculating VIF for 'enginesize'

```
X_train_eng= X_train_new.drop('enginesize',axis=1)
y_train_eng= X_train_new['enginesize']
X_train_eng_sm=sm.add_constant(X_train_eng)
model_eng=sm.OLS(y_train_eng,X_train_eng_sm).fit()
print(model_eng.summary())
```

OLS Regression Results

```
=====
Dep. Variable: enginesize R-squared: 0.79
Model: OLS Adj. R-squared: 0.79
Method: Least Squares F-statistic: 180.
Date: Sat, 07 Sep 2024 Prob (F-statistic): 8.94e-4
Time: 03:28:27 Log-Likelihood: -89.26
No. Observations: 143 AIC: 186.
Df Residuals: 139 BIC: 198.
Df Model: 3
Covariance Type: nonrobust
=====
      coef  std err      t    P>|t|   [0.025  0.975
-----
const    2.949e-17  0.038  7.7e-16  1.000  -0.076  0.07
curbweight  0.6446  0.063  10.248  0.000   0.520  0.76
horsepower  0.5234  0.068  7.647  0.000   0.388  0.65
citympg    0.2421  0.073  3.305  0.001   0.097  0.38
=====
Omnibus: 33.197 Durbin-Watson: 2.11
Prob(Omnibus): 0.000 Jarque-Bera (JB): 154.01
Skew: 0.672 Prob(JB): 3.61e-3
Kurtosis: 7.903 Cond. No. 3.7
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correct

```
print(f"VIF value for enginesize is {1/(1-0.8)}")
```

VIF value for enginesize is 5.00000000000001

```
# VIF for all the features
vif_df= pd.DataFrame()
vif_df['Features']= X_train_sm.columns
vif_df['VIF']= [variance_inflation_factor(X_train_sm.values, i) for i in ran
vif_df
```

```

/usr/local/lib/python3.10/dist-packages/statsmodels/regression/linear_model.p
    return 1 - self.ssr/self.centered_tss
/usr/local/lib/python3.10/dist-packages/statsmodels/stats/outliers_influence.
    vif = 1. / (1. - r_squared_i)
/usr/local/lib/python3.10/dist-packages/statsmodels/regression/linear_model.p
    return 1 - self.ssr/self.centered_tss
/usr/local/lib/python3.10/dist-packages/pandas/core/nanops.py:1010: RuntimeWa
    sqr = _ensure_numeric((avg - values) ** 2)

```

	Features	VIF
0	const	0.000000
1	symboling	7.083662
2	doornumber	4.732720
3	wheelbase	27.742502
4	carlength	28.311274
5	carwidth	15.679405
6	carheight	8.119666
7	curbweight	46.887888
8	cylindernumber	32.378272
9	enginesize	72.608025
10	boreratio	8.850441
11	stroke	4.959952
12	compressionratio	174.996708
13	horsepower	42.909032
14	peakrpm	6.599422
15	citympg	45.105971
16	highwaympg	36.622279
17	fueltype_gas	inf
18	aspiration_turbo	6.422393
19	carbody_hardtop	4.751051
20	carbody_hatchback	17.752973
21	carbody_sedan	24.107649
22	carbody_wagon	12.137629
23	drivewheel_fwd	20.113183
24	drivewheel_rwd	25.986411
25	enginelocation_rear	inf

Note: 'inf' and 'NaN' values may occur due to multicollinearity.

26 enginetype_dohcv 4.209954

27 enginetype_I inf

28 Understanding Hypothesis Testing

29 enginetype_ohc 10.221547
 enginetype_ohcf inf

From the summary report of the linear regression, you may observe that each feature variable has a **p-value** ($P > |t|$) associated with it. The p-value is one of the important statistics which can be used to eliminate features which are not relatively significant in our model.

30 fuelsystem_2bbl 24.247454
 fuelsystem_4bbl 2.909069

Hypothesis Testing

31 fuelsystem_idi inf
 fuelsystem_mfi NaN

Hypothesis Testing is basically testing an assumption that we make about a parameter. This assumption may or may not be true.

32 fuelsystem_mpfi 34.268449
 fuelsystem_spdi 7.400405

The steps followed in hypothesis testing are:

1. An initial assumption or hypothesis is made.
2. The validity of that hypothesis is tested.
3. If the hypothesis is found to be true, it is accepted otherwise it is rejected.

33 CarCompany_audi 8.873794
 CarCompany_bmw 9.677094

There are two types of hypothesis:

1. **Null hypothesis:** denoted by H_0 , is a general statement or an initial assumption which we make about a parameter.
2. **Alternative hypothesis:** denoted by H_1 or H_a , It is contrary to the null hypothesis. It is the hypothesis we would accept if our null hypothesis is found to be false.

34 CarCompany_buick 8.353640
 CarCompany_chevrolet 2.878173
 CarCompany_honda 17.971016
 CarCompany_isuzu 4.818289

In hypothesis testing, we need to gather enough evidence to either accept or reject our null hypothesis. There are two types of hypothesis tests that can be used for multiple linear regression:

- **F-test:** This test measures the overall significance of all the coefficients.
- **T-test:** This test measures the significance of each individual coefficient.

35 CarCompany_nissan 8.293474

51 F-test

52 CarCompany_peugeot inf
 CarCompany_plymouth 7.698632

The F-test is used to assess all the coefficients collectively. It validates whether any of the independent variables are significant.

53 CarCompany_porsche 10.565245
 CarCompany_renault 3.867446

The regression equation for the car price prediction model can be given as

54 CarCompany_saab 7.927021
 $\$Y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \dots + \beta_{59} x_{59} + \epsilon$

$\$Y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \dots + \beta_{59} x_{59} + \epsilon$

where, 55 CarCompany_subaru inf
 CarCompany_toyota 16.053616

56 CarCompany_volkswagen 11.491334

- x_1 is symboling
- x_2 is doorno

- x_3 is `wheelbase`

$\vdots \vdots \vdots$

- x_{59} is `wheelbase` and
- Y is the `price`

Step 1: Define null and alternative hypothesis

$H_0: \beta_1 = \beta_2 = \dots = \beta_{59} = 0$ i.e. all the regression coefficients are equal to zero.

$H_1: \beta_i \neq 0$, i.e. at least one of the coefficient is not zero.

- H_0 means that none of the feature or independent variables have a significant relationship with our target variable `price` and our model has no predictive capability.
- H_1 means that at least one feature variable has a significant relationship with our target variable `price`.

Step 2: Calculate the test statistic value (in case of F-test it is F-statistic value)

It is calculated as

$$F^* = \frac{\text{explained variance}}{\text{unexplained variance}} = \frac{\text{MSM}}{\text{MSE}}$$

where,

- MSM is the Mean of Squares for Model
- MSE is Mean of Squared Errors (or Residuals)

Further, MSM is calculated as

$$\text{MSM} = \frac{\text{SSM}}{\text{DFM}} = \frac{\sum(y_{\text{pred}} - \bar{y})^2}{p-1}$$

where,

- SSM is the Sum of Squares for Model
- DFM is Degrees of Freedom for Model
- p is the number of independent variables

Similarly, MSE is calculated as:

$$\text{MSE} = \frac{\text{SSE}}{\text{DFE}} = \frac{\sum(y - y_{\text{pred}})^2}{N-p}$$

where,

- SSE is the Sum of Squares for Errors
- DFE is Degrees of Freedom for Errors
- $\$N\$N\$$ is number of instances (or rows) in the dataset

Let's create `mean_sq_model()` and `mean_sq_error()` functions to calculate the MSM and MSE values using the above formulae respectively.

Note: You can also obtain the MSM and MSE values using the `mse_model` and `mse_resid` attributes respectively of `statsmodels.api` module.

```
#calculating f-stat
f_stat=model.mse_model/model.mse_resid
print(f"The F-statistics for the model {model.mse_model/model.mse_resid}")
```

The F-statistics for the model 67.50726851182478

```
#Calculating p-value
from scipy.stats import norm
p_val = 2*(1-norm.cdf(abs(f_stat)))
print(p_val)
```

0.0

`model.f_pvalue`

3.515694157936587e-53

Note: If p-value is below 0.05, the null hypothesis will be rejected.

The p-value that we obtained from F-test is equal to 0.00, so we can reject our null hypothesis and conclude that at least one of the independent variable has linear relationship with our target variable `price`. But, what is p-value?

What is meant by p-value?

The p-value is a probability value that helps us to determine whether our hypothesis is correct. The p-value for each feature tests the null hypothesis that there is no correlation between the feature and the target variable. Smaller the p-value, stronger is the evidence that you should reject null hypothesis. A p-value less than 0.05 is statistically significant. It indicates that there is less than 5% probability that the null hypothesis is correct. Therefore, we reject the null hypothesis, and accept the alternative hypothesis. However, a p-value greater than 0.05 indicates weak evidence and we fail to reject the null hypothesis.

The F-test for our model rejected the null hypothesis and concluded that at least one feature variable is significant and our model definitely possess predictive capability. Now, we will perform **t-test** to determine which variables are significant in predicting the price of a car and which are not.

▼ T-test

After concluding from the F-test that at least one feature variable is significant, now we may want to know which variables are significant. For this, we can do a **t-test** to find out which independent variable is making a useful contribution in the prediction of the dependent variable.

$$\text{Y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \dots + \beta_{59} x_{59} + \epsilon$$

where,

- x_1 is **symboling**
- x_2 is **doornumber**
- x_3 is **wheelbase**

\vdots

- x_{59} is **wheelbase** and
- Y is the **price**

For example, let us determine whether feature **symboling** is contributing significantly in the prediction of dependent variable **price**. We will follow the same steps as that of F-test.

Step 1: Define the null and alternative hypothesis

$H_0: \beta_1 = 0$ i.e. **symboling** and **price** are not linearly related

$H_1: \beta_1 \neq 0$ i.e. **symboling** and **price** are linearly related

Step 2: Calculate the test statistic value (in case of t-test, it is t-statistic value)

The t-statistic is calculated as:

$$t^* = \frac{\text{coefficient} - \text{hypothesized value}}{\text{standard error of coefficient}}$$

As the hypothesized value is usually 0, $t^* = \frac{\text{coefficient}}{\text{standard error of coefficient}}$

\textrm{standard error of coefficient}}\\$\\$

For our example above, the t-statistic is:

$\text{\$t*=\frac{\beta_1}{SE(\beta_1)}}\$$

The **standard error of coefficient (SE)** is an estimate of the standard deviation of the coefficient, the amount it varies across cases. Its formula is quite complicated.

However, we can obtain standard error for every coefficient by using `bse` attribute of `statsmodels.api` module. The `b` in `bse` stands for the coefficient β and `se` for standard errors.

```
bse_symboling= model.bse['symboling']
t1= model.params['symboling']/bse_symboling
print(t1)
```

-0.766449291655052

```
# p-value for t-stat
print(f"p-value for t-stat for symboling {2* (1-norm.cdf(abs(t1)))}")
```

p-value for t-stat for symboling 0.4434090131014836

Hence symboling and price are not linearly related hence symboling is insignificant.

```
# Check all the p_values
print(model.pvalues)

const          0.000002
symboling      0.445484
doornumber     0.356501
wheelbase       0.491750
carlength      0.356087
carwidth        0.003989
carheight       0.364582
curbweight      0.016007
cylindernumber 0.244234
enginesize      0.000297
boreratio       0.007926
stroke          0.789116
compressionratio 0.782639
horsepower      0.758939
peakrpm         0.012603
citympg          0.456021
highwaympg       0.755790
fueltype_gas    0.014296
aspiration_turbo 0.002748
carbody_hardtop 0.889279
carbody_hatchback 0.002640
carbody_sedan   0.005734
carbody_wagon    0.003059
drivewheel_fwd 0.461885
```

drivewheel_rwd	0.296761
enginelocation_rear	0.010376
enginetype_dohcv	0.925901
enginetype_l	0.192379
enginetype_ohc	0.725297
enginetype_ohcf	0.285501
enginetype_ohcv	0.288010
enginetype_rotor	0.008174
fuelsystem_2bb1	0.080563
fuelsystem_4bb1	0.663401
fuelsystem_idi	0.237300
fuelsystem_mfi	0.588545
fuelsystem_mpfi	0.293220
fuelsystem_spdi	0.502118
fuelsystem_spfi	0.413231
CarCompany_audi	0.994202
CarCompany_bmw	0.000376
CarCompany_buick	0.002628
CarCompany_chevrolet	0.078490
CarCompany_dodge	0.022523
CarCompany_honda	0.317686
CarCompany_isuzu	0.420871
CarCompany_jaguar	0.561791
CarCompany_mazda	0.104761
CarCompany_mercury	0.523097
CarCompany_mitsubishi	0.011527
CarCompany_nissan	0.043811
CarCompany_peugeot	0.192379
CarCompany_plymouth	0.013786
CarCompany_porsche	0.029307
CarCompany_renault	0.085115
CarCompany_saab	0.677703
CarCompany_subaru	0.016262

```
new_df = pd.DataFrame()
new_df['feature']=X_train.columns
new_df['pValues'] = model.pvalues.values[1:]
new_df=new_df[new_df['pValues']<=0.05]
new_df
```

	feature	pValues
4	carwidth	0.003989
6	curbweight	0.016007
8	enginesize	0.000297
9	boreratio	0.007926
13	peakrpm	0.012603
16	fueltype_gas	0.014296
17	aspiration_turbo	0.002748
19	carbody_hatchback	0.002640
20	carbody_sedan	0.005734
21	carbody_wagon	0.003059
24	enginelocation_rear	0.010376
30	enginetype_rotor	0.008174
39	CarCompany_bmw	0.000376
40	CarCompany_buick	0.002628
42	CarCompany_dodge	0.022523
48	CarCompany_mitsubishi	0.011527
49	CarCompany_nissan	0.043811
51	CarCompany_plymouth	0.013786
52	CarCompany_porsche	0.029307
55	CarCompany_subaru	0.016262

```
#rebuild the Linear regression model
X_new= X[new_df.feature]
X_train2, X_test2, y_train2, y_test2 = train_test_split(X_new, y, test_size

X_train_sm2 = sm.add_constant(X_train2)
model2 = sm.OLS(y_train2, X_train_sm2).fit()
print(model2.summary())
```

OLS Regression Results

Dep. Variable:	price	R-squared:	0.95
Model:	OLS	Adj. R-squared:	0.95
Method:	Least Squares	F-statistic:	132.
Date:	Sat, 07 Sep 2024	Prob (F-statistic):	3.34e-7
Time:	03:28:30	Log-Likelihood:	-1206.
No. Observations:	137	AIC:	2455
Df Residuals:	116	BIC:	2516
Df Model:	20		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025
<hr/>					
const	-5.285e+04	8791.716	-6.012	0.000	-7.03e+04
carwidth	737.6042	146.262	5.043	0.000	447.913
curbweight	4.3546	0.931	4.678	0.000	2.511
enginesize	75.7302	9.699	7.808	0.000	56.519
boreratio	-3150.6456	881.598	-3.574	0.001	-4896.761
peakrpm	1.4756	0.445	3.319	0.001	0.595
fuelytype_gas	1677.9896	756.809	2.217	0.029	179.035
aspiration_turbo	2363.7059	526.933	4.486	0.000	1320.049
carbody_hatchback	-3445.8579	829.557	-4.154	0.000	-5088.900
carbody_sedan	-2551.4401	806.653	-3.163	0.002	-4149.118
carbody_wagon	-4302.0295	944.400	-4.555	0.000	-6172.532
enginelocation_rear	6084.9218	2050.365	2.968	0.004	2023.916
enginetype_rotor	6127.0878	1403.152	4.367	0.000	3347.969
CarCompany_bmw	7540.5414	845.205	8.922	0.000	5866.507
CarCompany_buick	6300.5518	1136.994	5.541	0.000	4048.592
CarCompany_dodge	-922.0000	810.515	-1.138	0.258	-2527.327
CarCompany_mitsubishi	-2326.3139	687.578	-3.383	0.001	-3688.148
CarCompany_nissan	-1704.4872	647.541	-2.632	0.010	-2987.024
CarCompany_plymouth	-1600.7490	791.994	-2.021	0.046	-3169.394
CarCompany_porsche	6807.8825	1474.859	4.616	0.000	3886.738
CarCompany_subaru	820.4880	879.024	0.933	0.353	-920.530
<hr/>					
Omnibus:	1.936	Durbin-Watson:	2.05		
Prob(Omnibus):	0.380	Jarque-Bera (JB):	1.52		
Skew:	-0.118	Prob(JB):	0.46		
Kurtosis:	3.459	Cond. No.	3.41e+0		
<hr/>					

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correct
- [2] The condition number is large, 3.41e+05. This might indicate that there is a strong multicollinearity or other numerical problems.

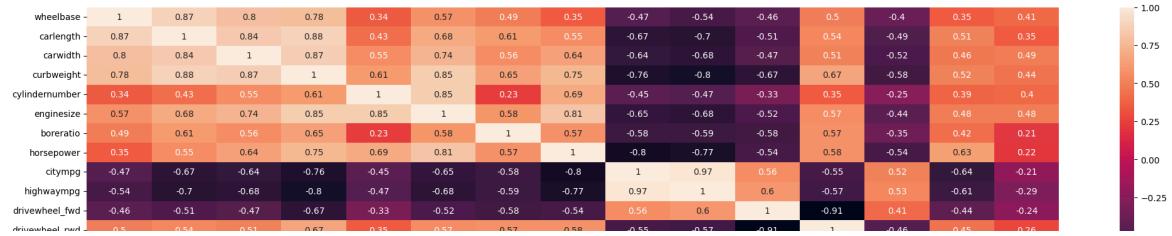
Still we have p_values higher than 0.05 in the result. A better approach to this is using RFE.

```
# Find the moderately too highly correlated features with price.
corr_df= df_new.corr()['price']
corr_data=corr_df[(corr_df >=0.5) | (corr_df<=-0.5)]
corr_features= list(corr_data.index)
corr_features.remove('price')
corr_features

['wheelbase',
 'carlength',
 'carwidth',
 'curbweight',
 'cylindernumber',
 'enginesize',
 'boreratio',
 'horsepower',
 'citympg',
```

```
'highwaympg',
'drivewheel_fwd',
'drivewheel_rwd',
'fuelsystem_2bb1',
'fuelsystem_mpfi',
'CarCompany_buick']
```

```
df2= df_new[corr_features]
plt.figure(figsize=(25,6),dpi=100)
sns.heatmap(df2.corr(),annot=True)
plt.show()
```



Recursive Feature Elimination (RFE)

Recursive feature elimination (RFE) is a feature selection (or elimination) method that fits a model and removes the weakest feature (or features). Here, you need to decide the numbers of features you want to select to build a model. Then you can validate your choice of number of features and increase or decrease them (if required).

Features are ranked by the model's `coef_` or `feature_importances_` attributes, and by recursively eliminating a small number of features per loop, RFE attempts to eliminate dependencies and collinearity that may exist in a machine learning model.

RFE requires a specified number of features to keep, however it is often not known in advance how many features are valid.

```
from sklearn.linear_model import LinearRegression
from sklearn.feature_selection import RFE
lr=LinearRegression()
rfe1= RFE(lr,n_features_to_select=10)
```

```
rfe1.fit(X_train[corr_features],y_train)
print(corr_features)
print(rfe1.support_)
print(rfe1.ranking_)

['wheelbase', 'carlength', 'carwidth', 'curbweight', 'cylindernumber', 'engin
[False False  True False  True  True  True  True  True  True False
 False  True  True]
[5 3 1 4 1 1 1 1 1 1 6 2 1 1]
```

```
# Look at the features selected.
rfe1_features= X_train[corr_features].columns[rfe1.support_]
rfe1_features

Index(['carwidth', 'cylindernumber', 'enginesize', 'boreratio',
'horsepower',
'citympg', 'highwaympg', 'drivewheel_fwd', 'fuelsystem_mpfi',
'CarCompany_buick'],
dtype='object')
```

```
# Check for multicollinearity.
X_train_rfe1= X_train[rfe1_features]
X_train_rfe1_sm = sm.add_constant(X_train_rfe1)
lr = sm.OLS(y_train, X_train_rfe1_sm).fit()
print(lr.summary())
```

OLS Regression Results						
=====						
Dep. Variable:	price	R-squared:	0.87			
Model:	OLS	Adj. R-squared:	0.86			
Method:	Least Squares	F-statistic:	91.0			
Date:	Sat, 07 Sep 2024	Prob (F-statistic):	3.16e-5			
Time:	03:28:39	Log-Likelihood:	-1336.			
No. Observations:	143	AIC:	2696			
Df Residuals:	132	BIC:	2728			
Df Model:	10					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	

const	1.412e+04	618.029	22.840	0.000	1.29e+04	1
carwidth	1730.4940	382.530	4.524	0.000	973.813	2
cylindernumber	-697.8044	630.166	-1.107	0.270	-1944.336	
enginesize	3477.6658	769.667	4.518	0.000	1955.187	5
boreratio	-700.5804	411.343	-1.703	0.091	-1514.258	
horsepower	2205.4249	595.424	3.704	0.000	1027.618	3
citympg	-683.5848	1078.676	-0.634	0.527	-2817.313	1
highwaympg	805.3504	1017.088	0.792	0.430	-1206.551	2
drivewheel_fwd	-2362.0657	716.205	-3.298	0.001	-3778.789	-
fuelsystem_mpfi	870.1314	695.994	1.250	0.213	-506.613	2
CarCompany_buick	8279.4418	1825.139	4.536	0.000	4669.136	1
=====						
Omnibus:	17.987	Durbin-Watson:	2.10			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	29.25			
Skew:	0.625	Prob(JB):	4.45e-0			
Kurtosis:	4.829	Cond. No.	17.			
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correct.

```
vif=pd.DataFrame()
vif['Feature']= X_train_rfe1_sm.columns
vif['VIF']=[variance_inflation_factor(X_train_rfe1_sm.values,i) for i in range(X_train_rfe1_sm.shape[1])]
vif
```

	Feature	VIF
0	const	6.532107
1	carwidth	2.502449
2	cylindernumber	6.791186
3	enginesize	10.130751
4	boreratio	2.893638
5	horsepower	6.062996
6	citympg	19.898364
7	highwaympg	17.691014
8	drivewheel_fwd	2.136322
9	fuelsystem_mpfi	2.070927
10	CarCompany_buick	1.548924

So 5 features have a higher p-value out of 10 selected features, so we need to remove 5 features. Rebuild the RFE model to select 5 features this time.

```
lr2=LinearRegression()
rfe2= RFE(lr2,n_features_to_select=5)
rfe2.fit(X_train[corr_features],y_train)
print(corr_features)
print(rfe2.support_)
print(rfe2.ranking_)

['wheelbase', 'carlength', 'carwidth', 'curbweight', 'cylindernumber', 'engin
[False False  True False False  True False  True False False  True False
 False False  True]
[10  8   1   9   3   1   4   1   6   5   1 11  7   2   1]
```

```
rfe2_features= X_train[corr_features].columns[rfe2.support_]
rfe2_features
```

```
Index(['carwidth', 'enginesize', 'horsepower', 'drivewheel_fwd',
       'CarCompany_buick'],
      dtype='object')
```

